

The Clean Web Economy: An In-Depth Design Exposition

Contents

1	Introduction	3
2	The Core Problem CWE Addresses	3
2.1	The Crisis of Trust and Attention	3
2.2	The Platform Monoculture	3
2.3	Why Not “Just Fix the Platforms”?	4
3	System Overview	4
3.1	Local Accounting	4
3.2	Zero-Knowledge Proof Submission	4
3.3	DAPR: The Payout Algorithm	5
3.4	Distributed Storage and Discovery	5
3.5	Identity and Governance	5
4	Common Misconceptions and Their Design Implications	5
4.1	“Isn’t This Just Another Subscription Service?”	5
4.2	“Does This Mean Creators Are Competing for Attention?”	5
4.3	“What Prevents the Richest Creators From Dominating?”	6
4.4	“Couldn’t Someone Fake Usage to Earn More?”	6
4.5	“Isn’t the System Too Complex for Ordinary Users?”	6
5	Architecture in Depth	6
5.1	The Client Layer	6
5.2	The Chain Layer	7
5.3	Discovery Layer	7
5.4	Storage Layer	7
5.5	DMF: Distributed Microservices Fabric	8
6	Economic Dynamics	8
6.1	The Tier Model	8
6.2	The “Free Rider” Problem	8
6.3	Sustainability Across Domains	8
7	Threat Model and Security Analysis	8
7.1	Sybil Resistance	8

7.2	Content Poisoning Attacks	9
7.3	Governance Capture	9
7.4	Legal Pressure Tactics	9
8	The Social Dynamics of a Non-Ad Web	9
8.1	Cultural Shifts	9
8.2	Community-Driven Curation	9
8.3	Global Equity	9
9	Limitations and Areas for Growth	10
9.1	Onboarding Barriers	10
9.2	Storage Redundancy	10
9.3	Proof Generation Costs	10
10	Conclusion	10

1 Introduction

The Clean Web Economy (CWE) began as a response to a pervasive unease about the state of the modern internet. By the mid-2020s, the dominant digital platforms—video streaming giants, social networks, news aggregators—had embraced a model based on profiling, targeted advertising, and algorithmic psychological steering. Even as public awareness grew, the underlying economic incentives remained unchanged: attention was currency, data was the commodity, and manipulation was the primary optimization target.

CWE proposed a radically different paradigm. Rather than removing ads from an existing platform, it rethought the economic structure from the ground up. Instead of extracting value from users, it redirected value to creators. Instead of centralizing trust, it minimized it through cryptography and decentralization. Instead of surveillance, it adopted privacy-first local computation.

Four years into its deployment, CWE has grown beyond a niche experiment. It now supports a diverse cultural ecosystem: small independent newsrooms, arthouse filmmakers, podcasters, musicians, scientific publishers, distributed microservices, educators, and researchers. Yet with growth comes scrutiny. Skepticism, both well-placed and ill-informed, deserves careful engagement.

This document offers a deep exploration of CWE’s design. It addresses doubts not as defensive rebuttals, but as opportunities to illuminate the architectural choices, constraints, and trade-offs that shape the system.

2 The Core Problem CWE Addresses

2.1 The Crisis of Trust and Attention

By the early 2020s, a consensus had formed that the mainstream web had become adversarial to its users. The core issues included:

- exploitative attention-harvesting algorithms,
- widespread behavioral profiling,
- misinformation amplified by engagement metrics,
- creator underpayment due to platform capture,
- opaque moderation and recommendation pipelines.

CWE’s thesis is not merely that ads are annoying; it is that the economic structure that incentivizes ads is incompatible with autonomy, creativity, and a healthy information ecosystem.

2.2 The Platform Monoculture

Centralization amplifies power asymmetries. Large platforms dictate:

- who gets visibility,
- which creators can sustain a career,

- what constitutes “quality” or “relevance”,
- how information flows through society.

CWE intentionally fragments infrastructure into modular components: fingerprint indexing, storage, identity verification, governance, and payment flows are all independent subsystems. No single actor can unilaterally control or capture the whole ecosystem.

2.3 Why Not “Just Fix the Platforms”?

Skeptics often ask whether existing platforms could adopt privacy-respecting, creator-friendly modes of operation. The answer is structurally no: their revenue depends on targeted advertising, and removing surveillance would eliminate their business model. Moreover, their centralized architectures fundamentally require trust in opaque private entities.

CWE replaces trust-based centralization with cryptographic verifiability and distributed governance.

3 System Overview

3.1 Local Accounting

Every user device maintains a local usage ledger. It records:

- content fingerprints,
- time spent,
- creator identifiers,
- optional voluntary metadata (never uploaded).

Crucially, the ledger never leaves the user’s device.

3.2 Zero-Knowledge Proof Submission

Once per cycle, the device sends a zero-knowledge proof to the chain layer. The proof states:

“I consumed content from creators A, B, and C in these proportions, and I am entitled to distribute my monthly tier accordingly.”

The chain verifies:

- the proof is well-formed,
- the proportions sum to 100%,
- the user holds a valid tier.

It learns nothing else.

3.3 DAPR: The Payout Algorithm

DAPR—Distributed Autonomous Proportional Reward—distributes funds from aggregated proofs to creators. It accounts for creator-set price floors, collaborative splits, and microservice dependencies.

Critically, DAPR does not reward engagement per se; it rewards verified usage proportionally weighted by value contributed.

3.4 Distributed Storage and Discovery

Content is stored in decentralized systems (IPFS, BitTorrent, or compatible layers). Discovery is handled by a separate subsystem that indexes fingerprints and metadata while using rate-limiting, reputation, and anti-spam techniques to resist capture.

3.5 Identity and Governance

Earning creators use verifiable credentials and decentralized identity systems. Governance uses jury-based arbitration, councils, and reproducible builds to minimize power concentration.

The result is a layered stack where no layer can unilaterally compromise user privacy or creator income.

4 Common Misconceptions and Their Design Implications

4.1 “Isn’t This Just Another Subscription Service?”

At first glance, CWE resembles a subscription platform. But the model differs fundamentally.

A traditional subscription ties money to a provider. A user pays Netflix, Spotify, or a news outlet. Their consumption is constrained by the platform.

CWE ties money to culture itself. Users contribute to a collective pool and direct its allocation proportionally to what they actually engage with. There is no “provider” in the usual sense; creators choose where to publish and users choose what to consume.

4.2 “Does This Mean Creators Are Competing for Attention?”

The system is intentionally designed to reduce the attention arms race.

- There is no recommendation algorithm optimizing for watch-time.
- There is no advantage to clickbait; engagement duration, not clicks, drives allocation.
- There is no infinite scroll.

Creators can focus on quality rather than virality. Indeed, several highly successful creators in CWE produce long-form essays, lectures, and films that thrive precisely because they appeal to shared interests rather than algorithmic exploitation.

4.3 “What Prevents the Richest Creators From Dominating?”

CWE addresses this through three mechanisms:

1. **No global visibility algorithm.** Discovery is decentralized; popularity does not create systemic bias.
2. **Price floors discourage over-commodification.** Creators cannot drop prices to predatory levels.
3. **Users are exposed to curated, community-led recommendation hubs.** No single entity controls them.

This creates a flat creative landscape where niche and minority creators can sustain themselves.

4.4 “Couldn’t Someone Fake Usage to Earn More?”

Cheating the system requires:

- forging content fingerprints,
- forging creator identities,
- forging local accounting entries,
- producing valid zero-knowledge proofs.

The architecture assumes adversarial users. Content fingerprints are collision-resistant. Local accounting is audited via challenge-response between client and chain. And the chain requires that fingerprints correspond to existing registered works with verifiable creators.

Fake content yields no earnings.

4.5 “Isn’t the System Too Complex for Ordinary Users?”

CWE’s complexity lives beneath the surface. For users, the experience is:

- install browser extension or player plugin,
- use the web normally,
- funds are distributed automatically.

Users need no knowledge of cryptography, distributed systems, or governance. The complexity is justified because it prevents exploitation and centralization.

5 Architecture in Depth

5.1 The Client Layer

The client software handles:

- fingerprint extraction,
- local usage tracking,

- ZK proof generation,
- privacy-preserving caching of metadata,
- optional sandboxing of proprietary players.

Fingerprinting is done via WASM A/V analysis for audio/video and cryptographic hashing for text and images.

5.2 The Chain Layer

The chain manages:

- tier verification,
- proof verification,
- DAPR payouts,
- splits between creators and collaborators,
- resource-backed token transactions (optional).

CWE runs on Ethereum Layer 2, using zk-rollups for scalability and privacy.

5.3 Discovery Layer

Discovery indexes fingerprints and metadata. Its design avoids:

- biased ranking,
- centralized moderation,
- recommendation capture.

Instead, communities maintain curated catalogs and tags.

5.4 Storage Layer

The storage layer is platform-neutral and supports:

- IPFS,
- BitTorrent,
- S3-compatible buckets,
- university archives,
- personal servers,
- local-first media caches.

Creators can self-host or use community nodes.

5.5 DMF: Distributed Microservices Fabric

DMF allows creators to define microservices:

- transcoding,
- filtering,
- analysis,
- moderation,
- recommendation engines,
- streaming gateways.

Each microservice is signed, discoverable, and can receive DAPR-based payment for usage.

6 Economic Dynamics

6.1 The Tier Model

The tier model replaces advertising revenue with a predictable, voluntary contribution system. Critics often assume users would be unwilling to pay. But empirical data from CWE shows that people already pay for multiple subscriptions; a unified model with direct creator payouts is more efficient.

6.2 The “Free Rider” Problem

CWE allows free browsing of public content. Non-paying users cannot generate proofs and thus contribute no economic value. Surprisingly, this does not destabilize the system, because a minority of paying users still support the ecosystem—similar to how Wikipedia functions with voluntary donations.

6.3 Sustainability Across Domains

- **Arthouse film:** long-form watch-time leads to sustainable flows.
- **News:** readership-based distribution avoids clickbait incentives.
- **Science:** readership micropayments replace paywalls and author fees.
- **Microservices:** operational costs are compensated through usage.

7 Threat Model and Security Analysis

7.1 Sybil Resistance

Tier membership requires cryptographic verification. Earnings require valid VC-based identity. Sybil attacks against earning identities are economically disincentivized because payouts require actual usage.

7.2 Content Poisoning Attacks

Fingerprinting and indexing resist poisoning via:

- multisource verification,
- rate-limited ingestion,
- reputation-based anti-spam filtering.

7.3 Governance Capture

The governance structure includes:

1. jury-based arbitration,
2. open councils with rotating membership,
3. reproducible code builds,
4. public decision logs,
5. no centralized administrative keys.

Capture requires compromising multiple independent systems simultaneously.

7.4 Legal Pressure Tactics

Because no central entity controls storage, indexing, or payouts, legal takedown attempts have limited effect. Creators remain individually responsible for their content; the infrastructure remains resilient.

8 The Social Dynamics of a Non-Ad Web

8.1 Cultural Shifts

CWE encourages slower, higher-quality media consumption. Without infinite scroll or algorithmic pressure, users explore based on interest rather than compulsion.

8.2 Community-Driven Curation

Curation shifts from platform heuristics to community lists, independent catalogs, and academic-style citation networks.

8.3 Global Equity

Creators in smaller markets benefit disproportionately, as the payout model is global and not tied to local ad inventory.

9 Limitations and Areas for Growth

9.1 Onboarding Barriers

Despite simplification, onboarding creators into SSI/VC systems remains challenging. UI/UX refinement and institutional support are ongoing needs.

9.2 Storage Redundancy

Distributed storage avoids centralization but increases complexity of ensuring availability. Community hosting incentives remain an active research area.

9.3 Proof Generation Costs

Local ZK proof generation, while optimized, remains computationally heavy on low-end devices. The roadmap includes GPU-accelerated clients and incremental proofs.

10 Conclusion

CWE is neither utopian nor naive. It is a pragmatic response to a dysfunctional web economy: a system built around verifiability, autonomy, and fairness. Its architecture is intentionally complex because the world it protects users from is adversarial. Its governance is intentionally democratic because centralization corrodes trust.

After four years of operation, CWE demonstrates that an alternative internet economy is not only possible but sustainable. It does not seek to replace the web with a new platform; it seeks to rebuild the web's economic foundations so that culture, knowledge, and public discourse can thrive without surveillance or manipulation.

The Clean Web Economy is still evolving, but its core principle remains constant: a web where value flows directly between users and creators, protected by privacy, strengthened by community, and governed by transparency.